

SARIMAX Forecasting of a Boarder Crossing Data Set

Inhaltsverzeichnis

1. Data set description..... 3

 Context..... 3

 Content..... 3

2. Data preprocessing..... 5

3. SARIMAX evaluation..... 9

1. Data set description

Context

The Bureau of Transportation Statistics (BTS) Border Crossing Data provide summary statistics for inbound crossings at the U.S.-Canada and the U.S.-Mexico border at the port level. Data are available for trucks, trains, containers, buses, personal vehicles, passengers, and pedestrians. Border crossing data are collected at ports of entry by U.S. Customs and Border Protection (CBP). The data reflect the number of vehicles, containers, passengers or pedestrians entering the United States. CBP does not collect comparable data on outbound crossings. Users seeking data on outbound counts may therefore want to review data from individual bridge operators, border state governments, or the Mexican and Canadian governments.

The data set is accessible at: <https://www.kaggle.com/datasets/akhilv11/border-crossing-entry-data>.

Content

COVERAGE: Incoming vehicle, container, passenger, and pedestrian counts at U.S.-Mexico and U.S.-Canada land border ports.

DEFINITIONS:

1. Bus Crossings: Number of arriving buses at a particular port, whether or not they are carrying passengers.
2. Container: A Container is defined as any conveyance entering the U.S. used for commercial purposes, either full or empty. Includes containers moving in-bond for the port initiating the bonded movements.
3. Types of Containers: The following are examples of a Container: Stakebed truck, truck with a car carrier, van, pickup truck/car, flatbed truck, piggyback truck with two linked trailers/containers = 2 containers, straight truck, bobtail truck, railcar, rail flatbed car stacked with four containers = 4 containers (on each rail car if there is multiple box containers count each container and the flatbed car.), and tri-level boxcar with multiple containers inside = 3 containers
4. Passengers Crossing In Buses: Number of persons arriving by bus requiring U.S. Customs and Border Protection (CBP) processing.
5. Passengers Crossing In Privately Owned Vehicles: Persons entering the United States at a particular port by private automobiles, pick-up trucks, motorcycles, recreational vehicles, taxis, ambulances, hearses, tractors, snowmobiles and other motorized private ground vehicles.
6. Passengers Crossing In Trains: Number of passengers and crew arriving by train and requiring CBP processing.

7. Pedestrian Crossings: The number of persons arriving on foot or by certain conveyance (such as bicycles, mopeds, or wheel chairs) requiring CBP processing.
8. Privately Owned Vehicle Crossings: Number of privately owned vehicles (POVs) arriving at a particular port. Includes pick-up trucks, motorcycles, recreational vehicles, taxis, snowmobiles, ambulances, hearses, and other motorized private ground vehicles.
9. Rail Container Crossings (loaded and empty): A container is any conveyance entering the U.S. used for commercial purposes, full or empty. In this case, it is the number of full or empty rail containers arriving at a port. This series includes containers moving as inbound shipments.
10. Train Crossings: Number of arriving trains at a particular port.
11. Truck Container Crossings (loaded and empty): A container is any conveyance entering the U.S. used for commercial purposes, full or empty. In this case, it is the number of full or empty truck containers arriving at a port. This series includes containers moving as inbound shipments.
12. Truck Crossings: Number of arriving trucks; does not include privately owned pick-up trucks.

Notes

Canada:

The ports of entry at Noyes, Minnesota and Whitetail, Montana closed in June 2006 and January 2013, respectively.

1. Incoming Trucks, Incoming PVs, PV Passengers, Incoming Buses, Bus Passengers, and Incoming Pedestrians

- o Bar Harbor and Portland, Maine (ferry crossing) - Ferries arrived from May to September. The Bar Harbor, Maine to Yarmouth, Nova Scotia ferry is no longer in operation.

- o Anacortes and Friday Harbor - The international ferries that connect Anacortes and Friday Harbor, Washington with Sidney, British Columbia do not run in February. Truck Containers (Loaded) and Truck Containers (Unloaded)

- o Passenger vehicle and passengers in personal vehicles data for Cape Vincent, New York (ferry) are available beginning in 2007. The ferry between Wolfe Island (Canada) and Cape Vincent does not operate in the winter.

2. Incoming Train Passengers

- o Includes both passengers and crew.

- o Starting with November 2017, Maine officials restrict international bridge traffic to passenger vehicles only.

Mexico:

Data for the port of Calexico are reported as a combined total with Calexico East.

1. Incoming Trucks:

o Data represent the number of truck crossings, not the number of unique vehicles, and include both loaded and unloaded trucks.

2. Incoming Train Passengers:

o Includes train crew. BTS is not aware of any passenger service currently operating across the U.S.-Mexico Border.

o CBP has indicated to BTS that since 2009 train crew are being exchanged at the Texas-Mexico border, and thus do not enter the United States.

2. Data preprocessing

Main question: Forecast the monthly number of boarder crossings between USA and Canada.

STEP1: We extract via DuckDB the relevant portion of the data set by means of the following list of SQL queries:

query1 = """

```
CREATE TABLE IF NOT EXISTS AllData AS
SELECT * FROM read_csv('Border_Crossing_Entry_Data.csv')
"""
```

Select the first 7 columns of the AllData table and store them into a new table "RelevantData"

query2 = """

```
CREATE TABLE IF NOT EXISTS RelevantData AS
SELECT #1, #2, #3, #4, #5, #6, #7
FROM AllData
"""
```

Select the first 7 columns of the "RelevantData" table and store them into a new table "US_Canada_Data"

query3 = """

```
CREATE TABLE IF NOT EXISTS US_Canada_Data AS
SELECT #1, #2, #3, #4, #5, #6, #7
FROM RelevantData
WHERE Border = 'US-Canada Border'
"""
```

Select the first 7 columns of the "RelevantData" table and store them into a new table "US_Mexico_Data"

query4 = """

```
CREATE TABLE IF NOT EXISTS US_Mexico_Data AS
SELECT #1, #2, #3, #4, #5, #6, #7
FROM RelevantData
WHERE Border = 'US-Mexico Border'
```

```

''''''
# Describe the schema of the US_Canada_Data table
query5 = "DESCRIBE US_Canada_Data"

# Convert a string date column to date format with a CTE
query6 = ''''''
CREATE TABLE IF NOT EXISTS US_Canada_Data_Altered AS
SELECT
    *,
    strptime(concat('01 ', Date), '%d %b %Y') AS converted_date
FROM
    US_Canada_Data
''''''

# Select the entire altered table US_Canada_Data_Altered
query7 = "SELECT * FROM US_Canada_Data_Altered"

# Describe the altered table US_Canada_Data_Altered
query8 = "DESCRIBE US_Canada_Data_Altered"

# Order the table US_Canada_Data_Altered based on the converted_date
# column in ascending manner
query9 = ''''''
SELECT *
FROM US_Canada_Data_Altered
ORDER BY converted_date ASC
''''''

# Create a US_Canada_AggregatedData table
query10 = ''''''
-- Step 1: Drop the table if it exists
DROP TABLE IF EXISTS US_Canada_AggregatedData;

-- Step 2: Create the new table and populate it with aggregated data
CREATE TABLE IF NOT EXISTS US_Canada_AggregatedData AS
SELECT
    converted_date,
    AVG(Value) AS avg_value_month,
    SUM(Value) AS total_value_month
FROM
    US_Canada_Data_Altered
GROUP BY
    converted_date
ORDER BY
    converted_date ASC;
''''''

# select the aggregated table:
query11 = ''''''
SELECT *

```

FROM US_Canada_AggregatedData

"""

These queries lead to the following table structure:

The aggregated time stamp ordered US_Canada_AggregatedData table:

converted_date timestamp	avg_value_month double	total_value_month int128
1996-01-01 00:00:00	10141.905982905982	9492824
1996-02-01 00:00:00	10653.472222222223	9971650
1996-03-01 00:00:00	11533.176282051281	10795053
1996-04-01 00:00:00	12576.271367521367	11771390
1996-05-01 00:00:00	14216.832264957266	13306955
1996-06-01 00:00:00	15745.551282051281	14737836
1996-07-01 00:00:00	17928.80876068376	16781365
1996-08-01 00:00:00	19217.96153846154	17988012
1996-09-01 00:00:00	14730.805555555555	13788034
1996-10-01 00:00:00	13659.291666666666	12785097
.	.	.
.	.	.
.	.	.
2024-03-01 00:00:00	12206.027675276753	6615667
2024-04-01 00:00:00	11444.026737967915	6420099
2024-05-01 00:00:00	13015.939609236235	7327974
2024-06-01 00:00:00	13756.374564459931	7896159
2024-07-01 00:00:00	16412.471304347826	9437171
2024-08-01 00:00:00	17274.277192982456	9846338
2024-09-01 00:00:00	12909.131715771231	7448569
2024-10-01 00:00:00	12736.114235500878	7246849
2024-11-01 00:00:00	11614.872072072072	6446254
2024-12-01 00:00:00	11762.11743772242	6610310
348 rows (20 shown)		3 columns

We convert the SQL table into a data frame via

```
„US_Canada_Data_Aggregated_df = db.sql("SELECT * FROM  
US_Canada_AggregatedData").df()",
```

leading to:

First ten entries of the US_Canada_AggregatedData_df:

	converted_date	avg_value_month	total_value_month
0	1996-01-01	10141.905983	9492824.0
1	1996-02-01	10653.472222	9971650.0
2	1996-03-01	11533.176282	10795053.0
3	1996-04-01	12576.271368	11771390.0
4	1996-05-01	14216.832265	13306955.0
5	1996-06-01	15745.551282	14737836.0
6	1996-07-01	17928.808761	16781365.0
7	1996-08-01	19217.961538	17988012.0
8	1996-09-01	14730.805556	13788034.0
9	1996-10-01	13659.291667	12785097.0

STEP 2: We create an aggregate dictionary from the extracted pandas DataFrame containing averaged numbers of boarder crossings for each month between January 1996 and December 2024 via:

```
AggregateDictionary = dict()
for iterYear in UniqueYears:
    MeasuresYear = []
    for iterMonth in UniqueMonths:
        condition1 = US_Canada_Data_df['Year'] == iterYear
        condition2 = US_Canada_Data_df['Month'] == iterMonth
        Temp_df = US_Canada_Data_df[condition1 & condition2]
        MeasuresYear.append(Temp_df['Value'].sum())
    AggregateDictionary[iterYear] = MeasuresYear

# sort dictionary according to its keys
myKeys = list(AggregateDictionary.keys())
myKeys.sort()

# Sorted Dictionary
sd = {i: AggregateDictionary[i] for i in myKeys}
AggregateDictionary = sd
print(AggregateDictionary)
```

This leads us to the following Data Frame structure:

C

```
AggregatedDF = pd.DataFrame.from_dict(AggregateDictionary)
print(AggregatedDF)
```

	1996	1997	1998	1999	2000	2001	2002	\
0	9492824	9436325	9509603	8963032	9120445	9968711	7967046	
1	11771390	10277874	11408750	10827952	11872809	10463283	9458774	
2	9971650	9533652	9606589	9515913	9417181	8803012	7839425	
3	10795053	11213563	10622465	10607585	11557273	10788933	9484491	
4	17988012	16834126	15867936	16408819	16509750	15909159	14371954	
5	13306955	12645355	12321480	14126950	13035044	11449120	10487686	
6	14737836	13397632	12780630	13600774	13653957	12655972	11102595	
7	10801524	10014597	9776730	10157514	9766290	7919578	8780240	
8	10819163	10903931	10270735	10842100	10567692	8136859	8850342	
9	12785097	12086771	11607099	12209136	12196150	8656311	10238366	
10	13788034	13345769	12836779	13211576	13369583	9942374	10848047	
11	16781365	15823115	15682994	16025977	16723499	15399079	13403729	
...								
	2003	2004	2005	...	2015	2016	2017	2018 \
0	8186226	7373490	7447288	...	6671807	6215665	6168767	6151874
1	7766297	8626328	8544541	...	7455737	6847059	7300505	7073413
2	7549888	7821507	7492452	...	6074264	5767485	5768250	5901475
3	8578623	8799951	8724428	...	7421462	6913898	6749117	7561801
4	12952628	12960841	12649601	...	10908827	10692018	10846607	10993925
5	9168311	9669448	9386894	...	8664906	8033685	7830360	8307472
6	10273715	10525294	10277161	...	9188601	8624151	8515363	9015529
7	8167744	8023030	8364507	...	6981127	6833933	7131453	7070154
8	8295872	8550079	8205193	...	6909179	6864575	6987074	6788073
9	9499606	9547455	9409051	...	7799362	7890198	7843460	7813441
10	9959044	10080151	10019526	...	8369654	8487012	8454019	8172138
11	12340460	12718048	12723059	...	11111636	10782099	10673299	10564665
...								
	2019	2020	2021	2022	2023	2024		
0	5841881	5895552	1596316	2126121	4685610	5235491		
1	7175976	1110625	1688595	4091393	5534787	6420099		
2	5451377	6000979	1465675	2140893	4536450	5553414		
3	7098470	3703717	1818274	2992770	5373629	6615667		
4	11146471	1714174	2333125	6550339	9156150	9846338		
5	8039300	1242037	1748625	4648165	6305656	7327974		
6	9762713	1540712	1825042	5273546	6978141	7896159		
7	7058135	1568631	2989342	4908379	6324285	6610310		
8	6763885	1624869	2674041	4999946	6059394	6446254		
9	7733413	1762684	2327083	5741597	6855550	7246849		
10	8070939	1730368	2312496	5388875	7226765	7448569		
11	10519310	1612690	1838438	6527388	9007752	9437171		

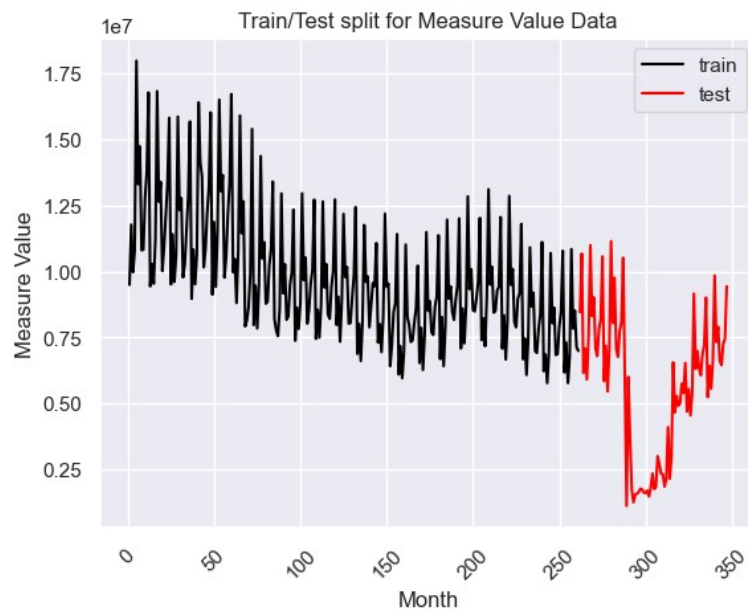
[12 rows x 29 columns]

STEP 3: We perform an 80-20 train-test split of the aggregated DataFrame via:


```
# create data set lists for training and testing
train = DataSet[:int(0.75*len(DataSet))]
test = DataSet[int(0.75*len(DataSet))+1:]

months1 = [(month+1) for month in range(len(train))]
months2 = [(month+1) for month in range(len(train),len(train)+len(test))]

plt.plot(months1,train, color = "black")
plt.plot(months2,test, color = "red")
plt.ylabel('Measure Value')
plt.xlabel('Month')
plt.xticks(rotation=45)
plt.title("Train/Test split for Measure Value Data")
plt.legend(['train', 'test'])
plt.show()
```



The test portion of the data set present a Covid-induced time anomaly in boarder crossings. This implies that we may have to expect our algorithm to perform suboptimally, however this split allows us to test the customized and automated forecasting performance of the cross-validated SARIMAX model under more „strained“ conditions.

3. SARIMAX evaluation

We trigger a customized Python function „SARIMAX_grid_search“ and obtain the following optimal SARIMAX parameters (seasonality s is set to 12) via

```
# start the grid search
start_time = time.time()
best_params = SARIMAX_grid_search(y, 12)
print("--- Duration: %s seconds ---" % (time.time() - start_time))

Tuned SARIMAX Parameters: (2, 1, 2, 2, 2, 2, 12)
Best average rmse score is 3622051.7420007405
```

	Time Period	True Values	Predicted Values	Confidence Lower \
0	0	7690676	7.876369e+06	4.669138e+06
1	1	7594110	7.503923e+06	4.259713e+06
2	2	8669030	7.609453e+06	4.243794e+06
3	3	9048099	8.223683e+06	4.826908e+06
4	4	11366681	8.724421e+06	5.278799e+06
5	5	6677975	1.023144e+07	6.749724e+06
6	6	8265991	6.450796e+06	2.929420e+06
7	7	6409080	8.117489e+06	4.558953e+06
8	8	7938922	6.341962e+06	2.745789e+06
9	9	11961504	7.836648e+06	4.203646e+06

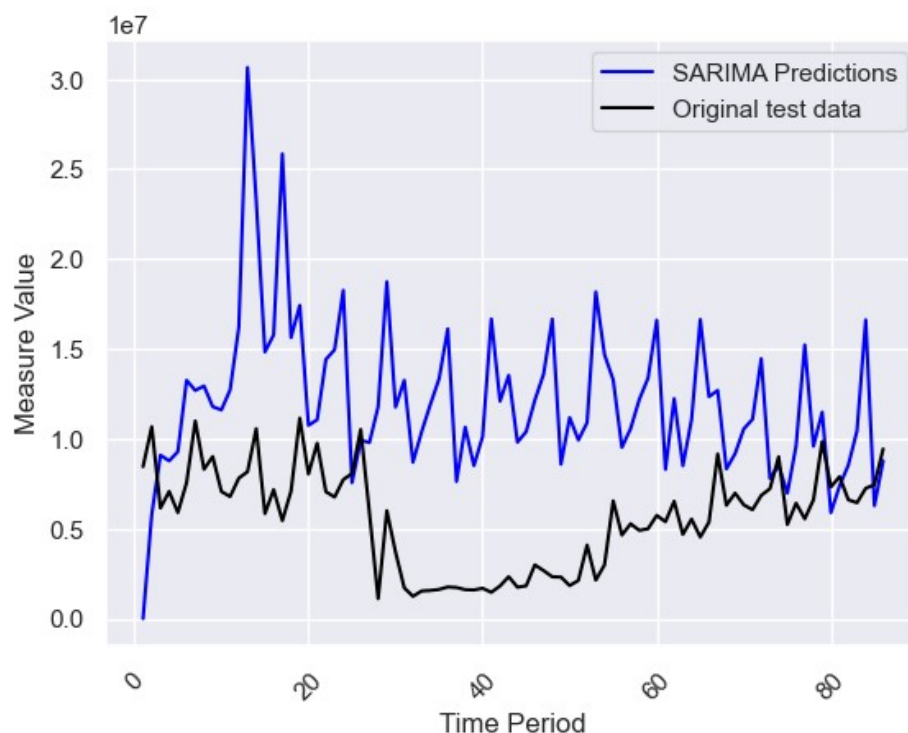
	Confidence Upper
0	1.108360e+07
1	1.074813e+07
2	1.097511e+07
3	1.162046e+07
4	1.217004e+07
5	1.371316e+07
6	9.972171e+06
7	1.167602e+07
8	9.938135e+06
9	1.146965e+07

```
--- Duration: 190.60557293891907 seconds ---

# print the optimal grid search parameters
print(best_params)


(2, 1, 2, 2, 2, 2, 12)
```

The average RMSE score indicates a lot of room for improvement, however due to the pronounced time series anomaly of the test portion within the extracted DataFrame it may be regarded as a solid „first guess“, as indicated by the corresponding forecast plot:



In the above plot „Measure Value“ denotes an aggregated number of monthly boarder crossings, whereas „Time Period“ refers to the number of months passed since the last training month.

Now we perform an automated cross-validated SARIMAX grid search by means of the Pythonic function „AutomatedSarimaxGridSearch“ and obtain the following results:



The ideal SARIMAX parameter set amounts to

``p=4.0`, `d=1.0`, `q=2.0`, `P=2.0`, `D=2.0`, `Q=2.0`, `s=12``

and leads to Mean Absolute Percentage Error (MAPE) of about 0.06 which indicates a significant forecasting improvement compared to the previous customized SARIMAX grid search.

Apparently, Pythonic SARIMAX modeling allows for robust and reliable trend predictions even in cases of highly anomalous time series trend patterns.